

Lean 4 for People Who Prove Things Elsewhere

A primer and cookbook for the mathematically fluent newcomer

Retrieved and assembled 2026-08-10

Abstract

This primer introduces the Lean 4 theorem prover and programming language to a reader with a strong mathematics and software-engineering background who has never used an *interactive theorem prover* (ITP)—a tool in which a human writes a proof step by step while the machine checks every move. We situate Lean against two systems such a reader may know, TLA^+ and Haskell; show through short, real code fragments what its two halves (a programming language and a proof assistant) actually look like; answer the practical questions of whether Lean is a general-purpose language and how it interoperates with C; and close with a worked, fully machine-checked theorem. Every Lean fragment in this document type-checks under Lean 4.33 with no external library, and lives in the companion `primer/` project.

Contents

1	What Lean 4 is	1
1.1	The functional-programming half, in code	2
1.2	What “theorem prover” means, concretely	2
2	Situating Lean: TLA^+ and Haskell	3
3	Is Lean a general-purpose language? IO, data, C interop	4
4	Tooling and editors	5
5	The proof experience and the crypto-relevant layer	5
6	Verifying properties of your own code	6
7	A worked theorem: the Josephus closed form	6
8	mathlib, and where to go next	7
	Glossary	8
	References	8

1 What Lean 4 is

Lean 4 is two things in one language: a **dependently-typed functional programming language** and an **interactive theorem prover**. The same syntax, type system, and tooling serve both jobs; you do not switch dialects to go from writing a program to writing a proof.

1.1 The functional-programming half, in code

“Functional programming language” means the ordinary things—definitions, pattern matching, recursion, algebraic data types—look like this:

```
def double (n : Nat) : Nat := n + n           -- a function Nat → Nat

def fib : Nat → Nat                         -- recursion with pattern matching
| 0 => 0
| 1 => 1
| n + 2 => fib n + fib (n + 1)

inductive Tree (α : Type) where             -- an algebraic data type
| leaf : Tree α
| node : Tree α → α → Tree α → Tree α

#eval fib 10                                -- 55 (#eval runs the code)
```

“Dependently-typed” is the part without a Haskell or ML analogue: a *type may mention a value*. The length of a vector can live *in its type*, so the type-checker—not a runtime check—enforces that appending an m -vector and an n -vector yields an $(m + n)$ -vector:

```
-- `Vector α n` is a list of α whose length n is part of the type.
example : Vector Nat 3 := #v[10, 20, 30]      -- ok
-- example : Vector Nat 3 := #v[10, 20]      -- TYPE ERROR: length 2 ≠ 3

-- The type of `append` promises the output length is the sum of the inputs':
#check (Vector.append : Vector α m → Vector α n → Vector α (m + n))
```

Because a type can depend on values, a type can also *encode a mathematical statement*—and that is precisely the bridge to theorem proving (Section 1.2). A second consequence, invisible above but load-bearing: Lean functions are **total by default**. The system checks that each recursion terminates (`fib` above is accepted because its arguments decrease); a function that might loop forever or crash is rejected unless you explicitly opt out. Totality is what makes it sound to read a function as a proof.

1.2 What “theorem prover” means, concretely

The link between the two halves is the **Curry–Howard correspondence**: a *proposition* is a *type*, and a *proof* of it is a *term* (a value) of that type. Proving P therefore means *constructing a term whose type is P* , and checking a proof is the same operation a compiler already performs—verifying that a term has the type it claims.

An aside on `rfl`. You will see `rfl` everywhere, so it is worth pinning down. It is the *one* canonical proof that a thing equals itself: its type says $a = a$ holds for every a . It proves a goal $x = y$ exactly when x and y *compute to the same value*—then, after evaluation, both sides are literally the same thing, and “ $a = a$ ” applies. So $2 + 2 = 4$ is closed by `rfl` because $2 + 2$ reduces to 4; but $a + b = b + a$ is *not*, because the two sides do not reduce to a common form (there you reach for `omega`).

A real proof, end to end. Something less trivial than $2 + 2$: the sum of the first n odd numbers is n^2 . The proof is a genuine induction with a real algebraic step.

```
def sumOdd : Nat → Nat
| 0 => 0
| n + 1 => sumOdd n + (2 * n + 1)  -- add the (n+1)-th odd number, 2n+1
```




Figure 1: How a proof is checked. Everything left of the dashed box may be elaborate or buggy; soundness rests only on the kernel re-checking the final term. Boxes are artifacts; solid arrows are transformations.

```

theorem sumOdd_eq (n : Nat) : sumOdd n = n * n := by
  induction n with
  | zero => rfl -- base: sumOdd 0 = 0 = 0*0, by computation
  | succ k ih =>
    have h : (k + 1) * (k + 1) = k * k + 2 * k + 1 := by
      simp only [Nat.succ_mul, Nat.mul_succ]; omega -- the one nonlinear fact
    simp only [sumOdd, ih, h] -- k*k + (2*k+1) = k*k + 2*k + 1
    omega
  
```

The two cases are the two ways to build a `Nat`: `zero` (closed by `rfl`, since both sides compute to 0) and `succ k`, where the *induction hypothesis* `ih : sumOdd k = k * k` is available and we extend it by one step.

What the system generates. A `by` block is a script; the proof is the *term* it produces, which is what the kernel checks. Take the two above.

You type:

```
theorem two_plus_two : 2 + 2 = 4 := by rfl
```

Lean generates (`#print two_plus_two`):

```
Eq.refl (2 + 2)
```

Why it proves the theorem: `Eq.refl` is the one constructor of equality, of type `a = a`. `Eq.refl (2 + 2)` therefore has type `2 + 2 = 2 + 2`—and since `2 + 2` and `4` reduce to the same value, that type *is* `2 + 2 = 4`. The kernel confirms the typing.

You type: the `sumOdd_eq` induction above. Lean generates (`#print sumOdd_eq`, plumbing elided):

```

fun n => Nat.recAux (Eq.refl (sumOdd 0)) -- the `zero` case
  (fun k ih => ...) -- the `succ` case, using ih
  n
  
```

Why it proves the theorem: `Nat.recAux` is induction on `Nat` as a function—give it a proof for 0 and a map from a proof for `k` to one for `k+1`, and it returns a proof for every `n`. Its type here is exactly `sumOdd n = n * n` for all `n`, which the kernel checks.

The kernel—a small, fixed core (a few thousand lines)—is the only thing trusted; tactics live outside it, so a buggy tactic cannot force a false theorem past it (Figure 1) [6]. A proof may also contain `sorry`, a hole that admits any goal; Lean warns while any remains.

2 Situating Lean: TLA⁺ and Haskell

The fastest way to understand Lean is by contrast with tools answering *different* questions (Table 1).

	TLA⁺	Lean 4	Haskell
Primary job	specify systems & temporal behavior	prove theorems; write verified programs	write programs
How you check	model-check a <i>finite</i> state space (TLC); finds counterexamples	<i>deduce</i> over infinite objects; no counterexamples unless you search	type-check; run tests
Types	untyped set theory	full <i>dependent</i> types; can express theorems	Hindley–Milner + extensions
Totality	—	total by default (termination-checked)	lazy & partial
Ask it	“bad interleaving?”	“true for all inputs, provably?”	“elegant program?”

Table 1: Three overlapping but distinct questions. Lean is the deductive, unbounded one.

Versus TLA⁺. TLA⁺ is a *specification language*: you model a system—often concurrent or distributed—as a state machine and state properties in temporal logic, then check them primarily by **model checking**, in which the TLC checker exhaustively explores a *finite* state space and prints a counterexample trace when a property fails. Lean does not explore states; it **proves** statements deductively over possibly-infinite objects. The mental shift is bounded-and-automatic-with-counterexamples (TLC) versus unbounded-and-deductive-without-counterexamples (Lean). Lean’s `#eval` and `decide` can feel faintly model-checker-like on finite decidable propositions, but that is evaluating a decision procedure, not exploring a temporal state space. Reach for TLA⁺ to ask “does my design have a bad interleaving?”; reach for Lean to ask “is this true for all inputs, provably?”.

Versus Haskell. Much will feel familiar: both are pure and functional with type classes, algebraic data types, `do` notation, monads, and strong inference. The differences are the point. Lean has full *dependent* types (Section 1.1); it is total by default where Haskell is lazy and partial (`undefined` inhabits every Haskell type); and its types can state theorems, which Haskell’s cannot in any practical sense. Treat Haskell as good intuition for the syntax and the functional style, then add “types can talk about values, and functions must terminate.”

3 Is Lean a general-purpose language? IO, data, C interop

Yes. Lean 4 is a genuine general-purpose language, not a proof-only toy. It compiles to C and then to native code, has an IO monad for real input/output, ships arrays, hash maps, strings and `ByteArray` for ordinary data manipulation, and is substantial enough that *the Lean compiler and much of its standard library are written in Lean itself* [6]. A minimal program:

```
def main : IO Unit := do
  let stdin ← IO.getStdin
  let line ← stdin.getLine
  IO.println s!"you said: {line.trim}"      -- string interpolation
```

Calling C from Lean. The `@[extern "sym"]` attribute binds a Lean declaration to an external C symbol; you give a Lean signature and type-level model, and link the C implementation:

```
@[extern "lean_sqrt"]      -- calls C function `double lean_sqrt(double)`
opaque cSqrt (x : Float) : Float
```

Calling Lean from C. The `@[export sym]` attribute exposes a Lean function under an unmangled, C-callable name; you compile the Lean code and link against the Lean runtime

(`libleanshared`), which provides garbage collection and the reference-counting discipline Lean’s compiled objects use:

```
@[export my_add]          -- becomes C-callable `lean_object* my_add(...)`
def myAdd (a b : Nat) : Nat := a + b
```

So interoperation runs both directions. The one caveat is the runtime: Lean-compiled code is C *plus* the Lean runtime and its object model (reference-counted `lean_object*`, with each parameter either owned or borrowed), so embedding Lean in a C program links a runtime rather than dropping in a freestanding `.o` [7]. Lean is thus usable for library building and data-processing tools; what it is *not* is a drop-in systems language with no runtime.

4 Tooling and editors

Lean 4 ships a **language server** (the `lean --server` command, speaking the *Language Server Protocol*, LSP), so editor support is a matter of an LSP client plus a plugin that renders the *infview*—the live display of the proof goal at the cursor.

- **VS Code** has the most complete, first-party extension. (Noted for completeness; a preference against it does not block anything below.)
- **Emacs** (`lean4-mode`) and **Neovim** (`lean.nvim`) are actively maintained LSP clients with *infview* support—the mainstream non-VS-Code choices.
- **Xcode: there is no Lean plugin.** Xcode does not expose a general third-party-language / LSP extension point, and no Lean 4 Xcode integration exists as of 2026. If you want to stay in the Apple toolchain, the practical route is the command line (below) plus an LSP-capable editor (Emacs/Neovim) alongside it.
- **Command line only** is entirely viable, and is how the companion project is checked:

```
lake build          -- type-check the whole project = check all
  proofs
lake env lean Primer/Basics.lean -- run one file, printing its #eval / #check
  output
```

There is no separate “test” step for the proofs: a clean `lake build` *is* the verification, because the build is exactly the kernel checking every term (Section 1.2). One practical caution learned while assembling this primer: `lake build` can report success from a cached build artifact; when you want the ground truth for a single file, run `lake env lean <file>`, which re-elaborates from source.

5 The proof experience and the crypto-relevant layer

Proofs are written either in *term mode* (hand Lean the term directly) or, more often, in *tactic mode*: a **by** block runs **tactics**—small commands that transform the goal while Lean assembles the underlying term. In an editor the *infview* shows the evolving goal; on the command line you reason it out. A working vocabulary: **rfl** (true by computation), **decide** (run a decision procedure on a finite/decidable goal), **simp** (rewrite with a lemma database), **omega** (linear arithmetic over integers and naturals), and **induction**.

Lean’s **BitVec** `n` type—fixed-width machine words with wraparound arithmetic—is the layer where cryptography lives, and it is directly relevant to the companion survey of SHA-256 formalization. SHA-256’s whole primitive layer is a few total functions:

```

def rotr (x : BitVec 32) (n : Nat) : BitVec 32 := (x >>> n) ||| (x <<< (32 - n))
def Ch    (x y z : BitVec 32) := (x &&& y) ^^^ ((~x) &&& z)
def Maj   (x y z : BitVec 32) := (x &&& y) ^^^ (x &&& z) ^^^ (y &&& z)
def Sigma0 (x : BitVec 32) := rotr x 2 ^^^ rotr x 13 ^^^ rotr x 22

-- Over ALL 256 eight-bit words, brute force (`decide`) settles this identity:
example : ∀ x : BitVec 8, x ^^^ x = 0#8 := by decide
-- Over 32 bits it is 232 cases; there `decide` is hopeless and the tactic is
-- `bv_decide` - bit-blast to SAT and check an LRAT certificate in-kernel.

```

That last tactic, `bv_decide`, is verified bit-blasting: it hands the goal to a SAT solver and then re-checks the solver’s refutation certificate with a Lean-verified checker, so the untrusted solver never enters the trusted base [3]. It is why Lean is now a strong platform for bit-level cryptographic reasoning even though a full software SHA-256 correctness proof in Lean does not yet exist—the subject of the companion survey.

6 Verifying properties of your own code

Set expectations precisely. Lean is excellent for proving properties of *pure functions*, data-structure *invariants*, and math-heavy *kernels* that you model or re-implement in Lean; there you can prove a statement for *all* inputs. Lean is *not* a push-button verifier for an arbitrary existing Rust, Python, or C repository—that is the domain of SMT-backed tools (*Satisfiability Modulo Theories*: automatic decision procedures over arithmetic, arrays, and the like) such as Dafny or Why3, or of model checkers such as TLA⁺. The productive loop is: model the critical pure core in Lean, prove it there, and treat that as your trustworthy reference. A representative shape—a fast implementation proved equal to an obvious reference model:

```

def sumTo : Nat → Nat                                     -- reference model
| 0 => 0
| n + 1 => (n + 1) + sumTo n

def sumToAcc : Nat → Nat → Nat                           -- optimized tail-recursive version
| 0, acc => acc
| n + 1, acc => sumToAcc n (acc + (n + 1))

theorem sumTo_correct (n : Nat) : sumToAcc n 0 = sumTo n := by
  simp [sumToAcc_eq]                                     -- via a generalized accumulator lemma

```

7 A worked theorem: the Josephus closed form

To show a non-trivial result carried all the way to a machine-checked proof, we formalize a classic from Knuth’s *Concrete Mathematics* [5]: the Josephus problem. Of n people in a circle every second is eliminated; the survivor’s position $J(n)$ satisfies

$$J(1) = 1, \quad J(2n) = 2J(n) - 1, \quad J(2n + 1) = 2J(n) + 1,$$

with the memorable closed form: writing $n = 2^m + \ell$ with $0 \leq \ell < 2^m$,

$$J(2^m + \ell) = 2\ell + 1$$

equivalently “ $J(n)$ is the binary numeral of n rotated left by one bit” (e.g. $13 = 1101_2 \mapsto 1011_2 = 11 = J(13)$). The bit-rotation reading is a direct callback to `rotr` in Section 5, which is why this problem is pleasantly close to the SHA-256 world. Here is the recurrence and the

general theorem, both machine-checked in pure Lean (no library). The listing below is **what you type**—the definition and a tactic proof:

```
def josephus : Nat → Nat
| 0 => 0
| 1 => 1
| n + 2 =>
  have : (n + 2) / 2 < n + 2 := Nat.div_lt_self (by omega) (by omega)
  if (n + 2) % 2 = 0 then 2 * josephus ((n + 2) / 2) - 1
  else 2 * josephus ((n + 2) / 2) + 1

theorem josephus_pow (m : Nat) :
  ∀ l, 1 < 2 ^ m → josephus (2 ^ m + 1) = 2 * l + 1 := by
  induction m with
  | zero => intro l hl; rw [Nat.pow_zero] at hl ⊢; have : l = 0 := by omega
    subst this; simp [josephus]
  | succ m ih =>
    intro l hl
    have hpow : (2 : Nat) ^ (m + 1) = 2 * 2 ^ m := by rw [Nat.pow_succ]; omega
    have hpos : 0 < (2 : Nat) ^ m := Nat.two_pow_pos m
    rcases (show 1 % 2 = 0 ∨ 1 % 2 = 1 by omega) with hpar | hpar
    · have hjm : 1 / 2 < 2 ^ m := by omega
      have hk : (2 : Nat) ^ (m+1) + 1 = 2 * (2 ^ m + 1 / 2) := by omega
      rw [hk, josephus_even (2 ^ m + 1 / 2) (by omega), ih (1 / 2) hjm]; omega
    · have hjm : 1 / 2 < 2 ^ m := by omega
      have hk : (2 : Nat) ^ (m+1) + 1 = 2 * (2 ^ m + 1 / 2) + 1 := by omega
      rw [hk, josephus_odd (2 ^ m + 1 / 2) (by omega), ih (1 / 2) hjm]; omega
```

Everything after `:= by` is a script of *tactics*—your instructions for building the proof, not the proof itself. **What Lean generates** from it is an ordinary term, which is what the kernel actually checks. Printing it (`#print josephus_pow`, trimmed here) shows the `induction` tactic became an application of `Nat`’s recursor, with your two cases as its arguments:

```
-- WHAT LEAN GENERATES (the proof term the kernel checks) - abbreviated:
theorem josephus_pow : ∀ (m l : Nat), 1 < 2 ^ m → josephus (2 ^ m + 1) = 2*l + 1 :=
fun m =>
  Nat.recAux
    (fun l hl => ...) -- the m = 0 case you wrote as `| zero => ...`
    (fun m ih l hl => ...) -- the m+1 case you wrote as `| succ m ih => ...`
    m
-- (the real term fills each `...` with ~15 lines of Eq.mpr/congrArg/... plumbing)
```

You never write that term by hand—that is the whole point of tactic mode. But it exists, and the kernel re-checks it; the tactics are only a convenient way to produce it. The proof itself is a textbook induction on m that splits on the parity of ℓ and steps down through the recurrence—exactly the *Concrete Mathematics* argument, now checked by the kernel rather than by a careful reader. The full file also machine-checks (over $1 \leq n \leq 2000$, via `native_decide`) that the recurrence, the closed form, and the one-bit-rotation form agree. This is the kind of concrete, self-contained result that Lean makes routine—and, notably, it required no mathematics library at all.

8 mathlib, and where to go next

mathlib is the community’s large, coherent library of formalized mathematics—algebra through analysis and beyond—with the lemmas and automation to match [1]. It is what you pull in when a proof rests on real mathematics; it is also a heavy dependency, which is why the examples here

deliberately use none of it (core Lean plus its standard library, *Batteries*, suffices for algorithmic and bit-level work). Good next steps: *Theorem Proving in Lean 4* [2] for the proof language, *Functional Programming in Lean* [4] for the programming side, and *Mathematics in Lean* [1] for working with mathlib. The companion **primer/** project holds every example above as runnable, type-checked source.

Glossary

Kernel

the small trusted core that checks a term has its claimed type; the sole component that must be trusted for soundness.

Term / Proposition

a term is a value/program; a proposition is a statement realized as a type whose terms are its proofs (Curry–Howard).

Tactic

a command that transforms the proof goal, building a proof term (e.g. **simp**, **omega**, **induction**).

Dependent type

a type that mentions a value (e.g. length-indexed vectors); what lets types state theorems.

ITP / ATP

interactive theorem prover (human-guided, machine-checked) versus automated theorem prover (search-driven).

LSP / FFI

Language Server Protocol (editor integration); Foreign Function Interface (calling between languages, here Lean↔C).

decide / bv_decide

run a decision procedure on a decidable goal; the latter bit-blasts **BitVec** goals to a SAT solver and re-checks its certificate.

elan / lake

the toolchain manager (installs/selects Lean versions) and the build system / package manager.

References

- [1] Jeremy Avigad and Patrick Massot. *Mathematics in Lean*. Local copy: `../refs/lean/mathematics_in_lean.pdf`, retrieved 2026-08-10. 2020. URL: https://leanprover-community.github.io/mathematics_in_lean/ (cit. on pp. 7, 8).
- [2] Jeremy Avigad et al. *Theorem Proving in Lean 4*. HTML only: https://leanprover.github.io/theorem_proving_in_lean4/, retrieved 2026-08-10. 2021. URL: https://leanprover.github.io/theorem_proving_in_lean4/ (cit. on p. 8).
- [3] Henrik Böving et al. “Interactive Bitvector Reasoning using Verified Bit-Blasting”. In: *Proc. ACM Program. Lang.* 9.OOPSLA2 (2025). No free PDF located (ACM Digital Library, gated), retrieved 2026-08-10, pp. 3259–3285. DOI: [10.1145/3763167](https://doi.org/10.1145/3763167). URL: <https://doi.org/10.1145/3763167> (cit. on p. 6).
- [4] David Thrane Christiansen. *Functional Programming in Lean*. HTML only: https://leanprover.github.io/functional_programming_in_lean/, retrieved 2026-08-10. 2023. URL: https://leanprover.github.io/functional_programming_in_lean/ (cit. on p. 8).

- [5] Ronald L. Graham, Donald E. Knuth, and Oren Patashnik. *Concrete Mathematics: A Foundation for Computer Science*. 2nd ed. Josephus problem, §1.3. Addison-Wesley, 1994 (cit. on p. 6).
- [6] Leonardo de Moura and Sebastian Ullrich. “The Lean 4 Theorem Prover and Programming Language”. In: *Automated Deduction – CADE 28*. Vol. 12699. Lecture Notes in Computer Science. Local copy: ../refs/lean/lean4-system-paper.pdf, retrieved 2026-08-10. Springer, 2021, pp. 625–635. DOI: 10.1007/978-3-030-79876-5_37. URL: <https://lean-lang.org/papers/lean4.pdf> (cit. on pp. 3, 4).
- [7] The Lean Community. *The Lean Language Reference: Foreign Function Interface*. Online reference manual (HTML). HTML only, retrieved 2026-08-10. 2026. URL: <https://lean-lang.org/doc/reference/latest/Run-Time-Code/Foreign-Function-Interface/> (cit. on p. 5).